

Calculating the Training Speed of an LLM Transformer Model in FLOPs

Estimating the number of floating-point operations (FLOPs) required to train a Transformer-based Large Language Model (LLM) is crucial for understanding the computational resources needed. This guide provides a detailed walkthrough on how to calculate the FLOPs for training the example model previously discussed.

Table of Contents

- [Calculating the Training Speed of an LLM Transformer Model in FLOPs](#)
 - [Table of Contents](#)
 - [Introduction](#)
 - [Why FLOPs Matter](#)
 - [Components Contributing to FLOPs](#)
 - [FLOPs Calculation per Component](#)
 - [Notation](#)
 - [Embedding Layer](#)
 - [Multi-Head Attention](#)
 - [Steps:](#)
 - [Calculations:](#)
 - [Feed-Forward Network](#)
 - [Layer Normalization](#)
 - [Activation Functions](#)
 - [Total FLOPs per Forward and Backward Pass](#)
 - [Forward Pass](#)
 - [Backward Pass](#)
 - [Total FLOPs per Training Step](#)
 - [Example Calculation](#)
 - [Given](#)
 - [Calculating FLOPs for One Layer](#)
 - [Embedding Layer FLOPs](#)
 - [Multi-Head Attention FLOPs](#)
 - [Feed-Forward Network FLOPs](#)
 - [Layer Normalization FLOPs](#)
 - [Total FLOPs per Layer](#)
 - [Total FLOPs for All Layers](#)
 - [Total FLOPs for Forward and Backward Pass](#)
 - [Additional Considerations](#)
 - [Conclusion](#)
 - [References](#)

Introduction

When training a neural network, especially large models like Transformers, it's important to estimate the computational cost. The number of floating-point operations (FLOPs) gives a hardware-agnostic measure of computational complexity.

Why FLOPs Matter

- **Resource Planning:** Helps in selecting appropriate hardware.
- **Performance Benchmarking:** Allows comparison between different models.
- **Energy Consumption Estimation:** Higher FLOPs often mean more energy usage.
- **Optimization:** Identifying bottlenecks to improve efficiency.

Components Contributing to FLOPs

The FLOPs count comes from:

1. **Embedding Layer**
2. **Multi-Head Attention (MHA)**
3. **Feed-Forward Network (FFN)**
4. **Layer Normalization**
5. **Activation Functions**

FLOPs Calculation per Component

Notation

- **Batch Size:** (B)
- **Sequence Length:** (L)
- **Model Dimension:** (d_{model})
- **Feed-Forward Dimension:** (d_{ff})
- **Number of Heads:** (h)
- **Number of Layers:** (N)

Embedding Layer

- **Operation:** Lookup + Addition
- **FLOPs per Token:**
 - Token Embedding: None (lookup operation)
 - Positional Embedding: None (lookup operation)
 - Addition of embeddings: (d_{model}) FLOPs per token
- **Total FLOPs:**
 - ($\text{Embedding Addition FLOPs}$) = $B \times L \times d_{\text{model}}$

Multi-Head Attention

Steps:

1. **Linear Projections:** Query (Q), Key (K), Value (V)

2. Scaled Dot-Product Attention

- Compute attention scores
- Apply softmax
- Compute attention output

3. Output Projection

Calculations:

1. Linear Projections (Q, K, V)

- **FLOPs per Projection:**
 - $(2 \times B \times L \times d_{\text{model}} \times d_{\text{model}})$
 - Factor of 2 accounts for multiply and add operations in matrix multiplication.
- **Total for Q, K, V:**
 - $(3 \times 2 \times B \times L \times d_{\text{model}}^2)$

2. Attention Scores

- **Compute Scores:**
 - $(B \times h \times L \times L \times \frac{d_{\text{k}}}{h})$
 - Since $(d_{\text{k}} = \frac{d_{\text{model}}}{h})$, this simplifies.
 - **FLOPs:**
 - $(2 \times B \times h \times L^2 \times \frac{d_{\text{model}}}{h}) = 2 \times B \times L^2 \times d_{\text{model}}$

3. Softmax

- **FLOPs:**
 - **Exponentials and Divisions:** Approximately (5) FLOPs per element.
 - Total FLOPs:
 - $(5 \times B \times h \times L^2)$

4. Weighted Sum

- **FLOPs:**
 - $(2 \times B \times h \times L^2 \times \frac{d_{\text{model}}}{h}) = 2 \times B \times L^2 \times d_{\text{model}}$

5. Output Projection

- **FLOPs:**
 - $(2 \times B \times L \times d_{\text{model}}^2)$

Feed-Forward Network

Consists of two linear transformations with an activation function in between.

1. First Linear Layer

- **FLOPs:**
 - $(2 \times B \times L \times d_{\text{model}} \times d_{\text{ff}})$

2. Activation Function

- **Assuming GELU Activation:**
 - Approximately (8) FLOPs per element.
 - **Total FLOPs:**
 - $(8 \times B \times L \times d_{\text{ff}})$

3. Second Linear Layer

- **FLOPs:**
 - $(2 \times B \times L \times d_{\text{ff}} \times d_{\text{model}})$

Layer Normalization

- **FLOPs per Layer Norm:**
 - **Mean Calculation:** $(B \times L \times d_{\text{model}})$ (sum and division)
 - **Variance Calculation:** $(B \times L \times d_{\text{model}})$ (subtract mean, square, sum, divide)
 - **Normalization:** $(B \times L \times d_{\text{model}})$ (subtract mean, divide by std)
 - **Total per Layer Norm:**
 - Approximately $(5 \times B \times L \times d_{\text{model}})$

Activation Functions

- **ReLU:** 1 FLOP per element
- **GELU:** Approximately 8 FLOPs per element

Total FLOPs per Forward and Backward Pass

Forward Pass

- Sum the FLOPs from all components for the forward pass.

Backward Pass

- The backward pass requires computing gradients, which often involves similar computations to the forward pass.
- **Rule of Thumb:** The backward pass takes approximately **2 to 3 times** the FLOPs of the forward pass.
- For estimation, we'll assume the backward pass is **2 times** the forward pass FLOPs.

Total FLOPs per Training Step

- **Total FLOPs:**
 - $(\text{FLOPs}_{\text{total}} = \text{FLOPs}_{\text{forward}} + \text{FLOPs}_{\text{backward}})$
 - $(\text{FLOPs}_{\text{backward}} = 2 \times \text{FLOPs}_{\text{forward}})$
 - Therefore, $(\text{FLOPs}_{\text{total}} = 3 \times \text{FLOPs}_{\text{forward}})$

Example Calculation

Given

- **Batch Size ((B))**: Let's assume (B = 1) for simplicity.
- **Sequence Length ((L))**: 1,024
- **Model Dimension ((d_{\text{model}}))**: 1,024
- **Feed-Forward Dimension ((d_{\text{ff}}))**: 4,096
- **Number of Heads ((h))**: 16
- **Number of Layers ((N))**: 24

Calculating FLOPs for One Layer

Embedding Layer FLOPs

- **Total FLOPs:**
 - (B \times L \times d_{\text{model}} = 1 \times 1,024 \times 1,024 = 1,048,576) FLOPs

Multi-Head Attention FLOPs

1. Linear Projections (Q, K, V)

- (3 \times 2 \times B \times L \times d_{\text{model}}^2)
- (6 \times 1 \times 1,024 \times (1,024)^2 = 6 \times 1,024 \times 1,048,576)
- **Result:** (6,442,450,944) FLOPs

2. Attention Scores

- (2 \times B \times L^2 \times d_{\text{model}})
- (2 \times 1 \times (1,024)^2 \times 1,024 = 2 \times 1,048,576 \times 1,024)
- **Result:** (2,147,483,648) FLOPs

3. Softmax

- (5 \times B \times h \times L^2)
- (5 \times 1 \times 16 \times (1,024)^2 = 5 \times 16 \times 1,048,576)
- **Result:** (83,886,080) FLOPs

4. Weighted Sum

- Same as Attention Scores: (2,147,483,648) FLOPs

5. Output Projection

- (2 \times B \times L \times d_{\text{model}}^2)
- (2 \times 1 \times 1,024 \times (1,024)^2 = 2,147,483,648) FLOPs

Total MHA FLOPs per Layer:

- Sum of the above: (6,442,450,944 + 2,147,483,648 + 83,886,080 + 2,147,483,648 + 2,147,483,648)
- **Result:** (12,968,788,968) FLOPs

Feed-Forward Network FLOPs

1. First Linear Layer

- $(2 \times B \times L \times d_{\text{model}} \times d_{\text{ff}})$
- $(2 \times 1 \times 1,024 \times 1,024 \times 4,096 = 8,589,934,592)$ FLOPs

2. GELU Activation

- $(8 \times B \times L \times d_{\text{ff}})$
- $(8 \times 1 \times 1,024 \times 4,096 = 33,554,432)$ FLOPs

3. Second Linear Layer

- $(2 \times B \times L \times d_{\text{ff}} \times d_{\text{model}})$
- Same as the first linear layer: $(8,589,934,592)$ FLOPs

Total FFN FLOPs per Layer:

- Sum of the above: $(8,589,934,592 + 33,554,432 + 8,589,934,592)$
- **Result:** $(17,213,423,616)$ FLOPs

Layer Normalization FLOPs

- Assuming two Layer Norms per layer:
 - $(2 \times 5 \times B \times L \times d_{\text{model}})$
 - $(10 \times 1 \times 1,024 \times 1,024 = 10,485,760)$ FLOPs

Total FLOPs per Layer

- **Sum:**
 - Embedding Layer: $(1,048,576)$ (only once)
 - MHA: $(12,968,788,968)$
 - FFN: $(17,213,423,616)$
 - Layer Norms: $(10,485,760)$
 - **Total per Layer:** $(30,192,698,344)$ FLOPs

Total FLOPs for All Layers

- **Total for All Layers:**
 - $(N \times \text{Total per Layer})$
 - $(24 \times 30,192,698,344 = 724,624,760,256)$ FLOPs
- **Add Embedding Layer FLOPs** (only once):
 - Total Forward FLOPs: $(724,624,760,256 + 1,048,576 \approx 724,625,808,832)$ FLOPs

Total FLOPs for Forward and Backward Pass

- **Forward Pass FLOPs:** $(\approx 724,625,808,832)$
- **Backward Pass FLOPs:** $(2 \times 724,625,808,832 = 1,449,251,617,664)$
- **Total FLOPs per Training Step:**

- $(724,625,808,832 + 1,449,251,617,664 = 2,173,877,426,496)$ FLOPs
-

Additional Considerations

1. Batch Size Impact:

- FLOPs scale linearly with batch size.
- For $(B = 16)$, total FLOPs would be $(16 \times)$ the calculated value.

2. Sequence Length Impact:

- FLOPs scale quadratically with sequence length in the attention mechanism.

3. Optimizations:

- **Sparse Attention:** Reduces FLOPs by focusing on local interactions.
- **Efficient Transformers:** Models designed to reduce computational complexity.

4. Hardware Efficiency:

- **Throughput:** Actual training speed depends on hardware (GPUs, TPUs) and their efficiency.
 - **Parallelism:** Utilizing data and model parallelism can affect training speed.
-

Conclusion

Calculating the FLOPs required to train a Transformer-based LLM involves summing the operations from embeddings, multi-head attention, feed-forward networks, and other components. For the example model with $(B = 1)$ and $(L = 1,024)$:

- **Total FLOPs per Training Step:** Approximately **2.17 trillion FLOPs**.

Understanding the computational requirements helps in selecting appropriate hardware and optimizing the training process.

References

- Vaswani, A., et al. (2017). ["Attention is All You Need"](#).
 - Kaplan, J., et al. (2020). ["Scaling Laws for Neural Language Models"](#).
 - Patterson, D., et al. (2021). ["Carbon Emissions and Large Neural Network Training"](#).
 - Narang, S., et al. (2021). ["Do Transformer Modifications Transfer Across Implementations and Applications?"](#).
-

Note: These calculations provide estimates. Actual FLOPs may vary based on implementation details, optimizations, and hardware-specific operations.