



# Calculating the Training Speed of an LLM Transformer Model in FLOPs

Estimating the number of floating-point operations (FLOPs) required to train a Transformer-based Large Language Model (LLM) is crucial for understanding the computational resources needed. This guide provides a detailed walkthrough on how to calculate the FLOPs for training the example model previously discussed.

---

## Table of Contents

- [Calculating the Training Speed of an LLM Transformer Model in FLOPs](#)
  - [Table of Contents](#)
  - [Introduction](#)
  - [Why FLOPs Matter](#)
  - [Components Contributing to FLOPs](#)
  - [FLOPs Calculation per Component](#)
    - [Notation](#)
    - [Embedding Layer](#)
    - [Multi-Head Attention](#)
      - [Steps:](#)
      - [Calculations:](#)
    - [Feed-Forward Network](#)
    - [Layer Normalization](#)
    - [Activation Functions](#)
  - [Total FLOPs per Forward and Backward Pass](#)
    - [Forward Pass](#)
    - [Backward Pass](#)
    - [Total FLOPs per Training Step](#)
  - [Example Calculation](#)
    - [Given](#)

- Calculating FLOPs for One Layer
    - Embedding Layer FLOPs
    - Multi-Head Attention FLOPs
    - Feed-Forward Network FLOPs
    - Layer Normalization FLOPs
    - Total FLOPs per Layer
  - Total FLOPs for All Layers
  - Total FLOPs for Forward and Backward Pass
  - Additional Considerations
  - Conclusion
  - References
- 

## Introduction

When training a neural network, especially large models like Transformers, it's important to estimate the computational cost. The number of floating-point operations (FLOPs) gives a hardware-agnostic measure of computational complexity.

---

## Why FLOPs Matter

- **Resource Planning:** Helps in selecting appropriate hardware.
  - **Performance Benchmarking:** Allows comparison between different models.
  - **Energy Consumption Estimation:** Higher FLOPs often mean more energy usage.
  - **Optimization:** Identifying bottlenecks to improve efficiency.
- 

## Components Contributing to FLOPs

The FLOPs count comes from:

1. **Embedding Layer**
  2. **Multi-Head Attention (MHA)**
  3. **Feed-Forward Network (FFN)**
  4. **Layer Normalization**
  5. **Activation Functions**
- 

# FLOPs Calculation per Component

## Notation

- **Batch Size:**  $B$
- **Sequence Length:**  $L$
- **Model Dimension:**  $d_{\text{model}}$
- **Feed-Forward Dimension:**  $d_{\text{ff}}$
- **Number of Heads:**  $h$
- **Number of Layers:**  $N$

## Embedding Layer

- **Operation:** Lookup + Addition
- **FLOPs per Token:**
  - Token Embedding: None (lookup operation)
  - Positional Embedding: None (lookup operation)
  - Addition of embeddings:  $d_{\text{model}}$  FLOPs per token
- **Total FLOPs:**
  - Embedding Addition FLOPs =  $B \times L \times d_{\text{model}}$

## Multi-Head Attention

### Steps:

1. **Linear Projections:** Query ( $Q$ ), Key ( $K$ ), Value ( $V$ )

## 2. Scaled Dot-Product Attention

- Compute attention scores
- Apply softmax
- Compute attention output

## 3. Output Projection

# Calculations:

### 1. Linear Projections (Q, K, V)

- **FLOPs per Projection:**

- $2 \times B \times L \times d_{\text{model}} \times d_{\text{model}}$
- Factor of 2 accounts for multiply and add operations in matrix multiplication.

- **Total for Q, K, V:**

- $3 \times 2 \times B \times L \times d_{\text{model}}^2$

### 2. Attention Scores

- **Compute Scores:**

- $B \times h \times L \times L \times \frac{d_k}{h}$
- Since  $d_k = \frac{d_{\text{model}}}{h}$ , this simplifies.
- **FLOPs:**

$$\blacksquare 2 \times B \times h \times L^2 \times \frac{d_{\text{model}}}{h} = 2 \times B \times L^2 \times d_{\text{model}}$$

### 3. Softmax

- **FLOPs:**

- **Exponentials and Divisions:** Approximately 5 FLOPs per element.
- Total FLOPs:

$$\blacksquare 5 \times B \times h \times L^2$$

### 4. Weighted Sum

- **FLOPs:**

- $2 \times B \times h \times L^2 \times \frac{d_{\text{model}}}{h} = 2 \times B \times L^2 \times d_{\text{model}}$

### 5. Output Projection

- **FLOPs:**

- $2 \times B \times L \times d_{\text{model}}^2$

# Feed-Forward Network

Consists of two linear transformations with an activation function in between.

### 1. First Linear Layer

- **FLOPs:**

- $2 \times B \times L \times d_{\text{model}} \times d_{\text{ff}}$

### 2. Activation Function

- **Assuming GELU Activation:**

- Approximately 8 FLOPs per element.

- **Total FLOPs:**

- $8 \times B \times L \times d_{\text{ff}}$

### 3. Second Linear Layer

- **FLOPs:**

- $2 \times B \times L \times d_{\text{ff}} \times d_{\text{model}}$

## Layer Normalization

- **FLOPs per Layer Norm:**

- **Mean Calculation:**  $B \times L \times d_{\text{model}}$  (sum and division)

- **Variance Calculation:**  $B \times L \times d_{\text{model}}$  (subtract mean, square, sum, divide)

- **Normalization:**  $B \times L \times d_{\text{model}}$  (subtract mean, divide by std)

- **Total per Layer Norm:**

- Approximately  $5 \times B \times L \times d_{\text{model}}$

## Activation Functions

- **ReLU:** 1 FLOP per element

- **GELU:** Approximately 8 FLOPs per element

---

## Total FLOPs per Forward and Backward Pass

### Forward Pass

- Sum the FLOPs from all components for the forward pass.

## Backward Pass

- The backward pass requires computing gradients, which often involves similar computations to the forward pass.
- **Rule of Thumb:** The backward pass takes approximately **2 to 3 times** the FLOPs of the forward pass.
- For estimation, we'll assume the backward pass is **2 times** the forward pass FLOPs.

## Total FLOPs per Training Step

- **Total FLOPs:**
    - $\text{FLOPs}_{\text{total}} = \text{FLOPs}_{\text{forward}} + \text{FLOPs}_{\text{backward}}$
    - $\text{FLOPs}_{\text{backward}} = 2 \times \text{FLOPs}_{\text{forward}}$
    - Therefore,  $\text{FLOPs}_{\text{total}} = 3 \times \text{FLOPs}_{\text{forward}}$
- 

## Example Calculation

### Given

- **Batch Size ( $B$ ):** Let's assume  $B = 1$  for simplicity.
- **Sequence Length ( $L$ ):** 1,024
- **Model Dimension ( $d_{\text{model}}$ ):** 1,024
- **Feed-Forward Dimension ( $d_{\text{ff}}$ ):** 4,096
- **Number of Heads ( $h$ ):** 16
- **Number of Layers ( $N$ ):** 24

## Calculating FLOPs for One Layer

### Embedding Layer FLOPs

- **Total FLOPs:**
  - $B \times L \times d_{\text{model}} = 1 \times 1,024 \times 1,024 = 1,048,576$  FLOPs

## Multi-Head Attention FLOPs

### 1. Linear Projections (Q, K, V)

- $3 \times 2 \times B \times L \times d_{\text{model}}^2$
- $6 \times 1 \times 1,024 \times (1,024)^2 = 6 \times 1,024 \times 1,048,576$
- **Result:** 6,442,450,944 FLOPs

### 2. Attention Scores

- $2 \times B \times L^2 \times d_{\text{model}}$
- $2 \times 1 \times (1,024)^2 \times 1,024 = 2 \times 1,048,576 \times 1,024$
- **Result:** 2,147,483,648 FLOPs

### 3. Softmax

- $5 \times B \times h \times L^2$
- $5 \times 1 \times 16 \times (1,024)^2 = 5 \times 16 \times 1,048,576$
- **Result:** 83,886,080 FLOPs

### 4. Weighted Sum

- Same as Attention Scores: 2,147,483,648 FLOPs

### 5. Output Projection

- $2 \times B \times L \times d_{\text{model}}^2$
- $2 \times 1 \times 1,024 \times (1,024)^2 = 2,147,483,648$  FLOPs

## Total MHA FLOPs per Layer:

- Sum of the above:  $6,442,450,944 + 2,147,483,648 + 83,886,080 + 2,147,483,648 + 2,147,483,648$
- **Result:** 12,968,788,968 FLOPs

## Feed-Forward Network FLOPs

### 1. First Linear Layer

- $2 \times B \times L \times d_{\text{model}} \times d_{\text{ff}}$
- $2 \times 1 \times 1,024 \times 1,024 \times 4,096 = 8,589,934,592$  FLOPs

### 2. GELU Activation

- $8 \times B \times L \times d_{\text{ff}}$
- $8 \times 1 \times 1,024 \times 4,096 = 33,554,432$  FLOPs

### 3. Second Linear Layer

- $2 \times B \times L \times d_{\text{ff}} \times d_{\text{model}}$
- Same as the first linear layer: 8,589,934,592 FLOPs

### Total FFN FLOPs per Layer:

- Sum of the above:  $8,589,934,592 + 33,554,432 + 8,589,934,592$
- **Result:** 17,213,423,616 FLOPs

### Layer Normalization FLOPs

- Assuming two Layer Norms per layer:
  - $2 \times 5 \times B \times L \times d_{\text{model}}$
  - $10 \times 1 \times 1,024 \times 1,024 = 10,485,760$  FLOPs

### Total FLOPs per Layer

- **Sum:**
  - Embedding Layer: 1,048,576 (only once)
  - MHA: 12,968,788,968
  - FFN: 17,213,423,616
  - Layer Norms: 10,485,760
  - **Total per Layer:** 30,192,698,344 FLOPs

### Total FLOPs for All Layers

- **Total for All Layers:**
  - $N \times \text{Total per Layer}$
  - $24 \times 30,192,698,344 = 724,624,760,256$  FLOPs
- **Add Embedding Layer FLOPs** (only once):
  - Total Forward FLOPs:  $724,624,760,256 + 1,048,576 \approx 724,625,808,832$  FLOPs

### Total FLOPs for Forward and Backward Pass

- **Forward Pass FLOPs:**  $\approx 724,625,808,832$
  - **Backward Pass FLOPs:**  $2 \times 724,625,808,832 = 1,449,251,617,664$
  - **Total FLOPs per Training Step:**
    - $724,625,808,832 + 1,449,251,617,664 = 2,173,877,426,496$  FLOPs
-



# Additional Considerations

## 1. Batch Size Impact:

- FLOPs scale linearly with batch size.
- For  $B = 16$ , total FLOPs would be  $16\times$  the calculated value.

## 2. Sequence Length Impact:

- FLOPs scale quadratically with sequence length in the attention mechanism.

## 3. Optimizations:

- **Sparse Attention:** Reduces FLOPs by focusing on local interactions.
- **Efficient Transformers:** Models designed to reduce computational complexity.

## 4. Hardware Efficiency:

- **Throughput:** Actual training speed depends on hardware (GPUs, TPUs) and their efficiency.
  - **Parallelism:** Utilizing data and model parallelism can affect training speed.
- 

# Conclusion

Calculating the FLOPs required to train a Transformer-based LLM involves summing the operations from embeddings, multi-head attention, feed-forward networks, and other components. For the example model with  $B = 1$  and  $L = 1,024$ :

- **Total FLOPs per Training Step:** Approximately **2.17 trillion FLOPs**.

Understanding the computational requirements helps in selecting appropriate hardware and optimizing the training process.

---

# References

- Vaswani, A., et al. (2017). ["Attention is All You Need"](#).

- Kaplan, J., et al. (2020). ["Scaling Laws for Neural Language Models"](#).
  - Patterson, D., et al. (2021). ["Carbon Emissions and Large Neural Network Training"](#).
  - Narang, S., et al. (2021). ["Do Transformer Modifications Transfer Across Implementations and Applications?"](#).
- 

*Note: These calculations provide estimates. Actual FLOPs may vary based on implementation details, optimizations, and hardware-specific operations.*